

Machine Learning Programming Assignment Week 4

Rutger van der Meer

2024-04-05

Prepare Data

```
# Load necessary libraries  
library(caret)
```

```
## Loading required package: ggplot2
```

```
## Loading required package: lattice
```

```
library(randomForest)
```

```
## randomForest 4.7-1.1
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##  
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:ggplot2':  
##  
## margin
```

```

training.url <- 'http://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv'
test.cases.url <- 'http://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv'

# Define a function named 'downloadcsv' that takes two parameters:
# 1. 'url': a character string representing the URL from which the CSV file will be downloaded
# 2. 'nastrings': a character vector containing strings that should be treated as missing values

downloadcsv <- function(url, nastrings) {

  # Create a temporary file name and assign it to the 'temp' variable
  temp <- tempfile()

  # Download the CSV file from the specified URL and save it to the temporary file using the 'curl' method
  download.file(url, temp, method = "curl")

  # Read the CSV file into R as a dataframe, specifying 'na.strings' to handle missing values
  data <- read.csv(temp, na.strings = nastrings)

  # Remove the temporary file from the system
  unlink(temp)

  # Return the dataframe containing the downloaded CSV data
  return(data)
}

# Download the training data CSV file from the specified URL using the 'downloadcsv' function,
# treating empty strings, "NA", and "#DIV/0!" as missing values
train <- downloadcsv(training.url, c("", "NA", "#DIV/0!"))
test <- downloadcsv(test.cases.url, c("", "NA", "#DIV/0!"))

# Set the random seed for reproducibility
set.seed(123456)

# Create a data partition for the training data, with 80% of the data used for training and the rest for validation
trainset <- createDataPartition(train$classe, p = 0.8, list = FALSE)

# Subset the training data using the generated partition, creating a training dataset ('Training') and a validation dataset ('Validation')
Training <- train[trainset, ]
Validation <- train[-trainset, ]

```

Feature selection

```
# Check for near zero variance predictors in the training dataset and drop them if necessary
# The variable 'nonzerocol' will store the indices of near zero variance predictors
nonzerocol <- nearZeroVar(Training)

# Remove the near zero variance predictors from the training dataset
Training <- Training[, -nonzerocol]

# exclude columns with 40% more missing values exclude descriptive columns

countlength <- sapply(Training, function(x) {
  sum(!(is.na(x) | x == ""))
})

nullCol <- names(countlength[countlength < 0.6 * length(Training$classe)])

descriptcol <- c("X", "user_name", "raw_timestamp_part_1", "raw_timestamp_part_2",
  "cvtd_timestamp", "new_window", "num_window")

excludecolumns <- c(descriptcol, nullCol)

Training <- Training[, !names(Training) %in% excludecolumns]
```

Build Training Set

```
# Build a random forest model using the training dataset ('Training')
# The response variable 'classe' is converted to a factor using 'as.factor' for classification
# All other variables are considered as predictors ('.')
# The model is configured to calculate variable importance and consists of 10 trees ('ntrees = 10')
rfModel <- randomForest(as.factor(classe) ~ ., data = Training, importance = TRUE, ntrees = 10)

## Model Validation

# Generate predictions on the training dataset using the trained random forest model
ptraining <- predict(rfModel, Training)

# Use 'union' to ensure the same levels are used for both predicted and actual values
u1 <- union(ptraining, Training$classe)

# Create a confusion matrix comparing the predicted and actual values, using the 'table' function
# 'factor' is used to ensure that both predicted and actual values have the same levels ('u1')
t1 <- table(factor(ptraining, u1), factor(Training$classe, u1))

# Print the confusion matrix as a formatted output using the 'confusionMatrix' function
print(confusionMatrix(t1))
```

```
## Confusion Matrix and Statistics
##
##
##      A      B      C      D      E
## A 4464      0      0      0      0
## B      0 3038      0      0      0
## C      0      0 2738      0      0
## D      0      0      0 2573      0
## E      0      0      0      0 2886
##
## Overall Statistics
##
##              Accuracy : 1
##              95% CI : (0.9998, 1)
##      No Information Rate : 0.2843
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 1
##
##      McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##              Class: A Class: B Class: C Class: D Class: E
## Sensitivity          1.0000   1.0000   1.0000   1.0000   1.0000
## Specificity          1.0000   1.0000   1.0000   1.0000   1.0000
## Pos Pred Value       1.0000   1.0000   1.0000   1.0000   1.0000
## Neg Pred Value       1.0000   1.0000   1.0000   1.0000   1.0000
## Prevalence           0.2843   0.1935   0.1744   0.1639   0.1838
## Detection Rate       0.2843   0.1935   0.1744   0.1639   0.1838
## Detection Prevalence 0.2843   0.1935   0.1744   0.1639   0.1838
## Balanced Accuracy    1.0000   1.0000   1.0000   1.0000   1.0000
```

Our model performs good against training set, But we will cross validate against the held out set and check if we have avoided overfitting.

Cross validation

```
# Generate predictions on the validation dataset using the trained random forest model
pvalidation <- predict(rfModel, Validation)

# Use 'union' to ensure the same levels are used for both predicted and actual values
u2 <- union(pvalidation, Validation$classe)

# Create a confusion matrix comparing the predicted and actual values on the validation dataset,
# using the 'table' function
# 'factor' is used to ensure that both predicted and actual values have the same levels ('u2')
t2 <- table(factor(pvalidation, u2), factor(Validation$classe, u2))

# Print the confusion matrix as a formatted output using the 'confusionMatrix' function
print(confusionMatrix(t2))
```

```
## Confusion Matrix and Statistics
##
##
##      A      B      C      D      E
## A 1116      1      0      0      0
## B      0    758      0      0      0
## C      0      0   684      4      0
## D      0      0      0   638      3
## E      0      0      0      1   718
##
## Overall Statistics
##
##              Accuracy : 0.9977
##              95% CI : (0.9956, 0.999)
##      No Information Rate : 0.2845
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.9971
##
##      McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##              Class: A Class: B Class: C Class: D Class: E
## Sensitivity          1.0000   0.9987   1.0000   0.9922   0.9958
## Specificity          0.9996   1.0000   0.9988   0.9991   0.9997
## Pos Pred Value       0.9991   1.0000   0.9942   0.9953   0.9986
## Neg Pred Value       1.0000   0.9997   1.0000   0.9985   0.9991
## Prevalence           0.2845   0.1935   0.1744   0.1639   0.1838
## Detection Rate       0.2845   0.1932   0.1744   0.1626   0.1830
## Detection Prevalence 0.2847   0.1932   0.1754   0.1634   0.1833
## Balanced Accuracy    0.9998   0.9993   0.9994   0.9957   0.9978
```

Cross validation accuracy is 99.7% & out-of-sample error is 0.3%. So the model performs excellent

Test set prediction

Prediction of our algorithm for the test set is

```
ptest <- predict(rfModel, test)
ptest
```

```
##  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20
## B  A  B  A  A  E  D  B  A  A  B  C  B  A  E  E  A  B  B  B
## Levels: A B C D E
```